

1. `grep -w` checks for whole words, not just substrings
2. `grep -o` will only return the matches (not the entire line)
3. The `-` character also has to be escaped
4. `NR` = number of total lines; `NF` = number of fields in the current line
5. `length($1)`
6. `head -n 10 file.txt` => prints first 10 lines
7. `tail -n 10` = prints last 10 lines
8. `tail --lines=+2 newfile.txt` - starts from the 2nd line (1-indexed)
9. `sed -i` will actually replace things in the file
10. `awk -F[:] 'BEGIN {sum = 0} $13 > $14 {sum += $13} END {print sum}' last.fake`
11. `\>` means end of word
 - a. `$` means end of line
 - b. These 2 can be combined
12. we can only access the first 9 arguments; we need "shift n" to get the next ones
13. `lineCount=`cat file | wc -l``
14. `wc -m` = number of characters
 - a. `wc -w` = number of words
15. `cut -c 2-4 file.txt` = only displays the characters between 2 and 4 from `file.txt`
16. `x=`expr $x + 1`` ⇔ `x++`
17. checks if it's even: ``expr $x % 2` -eq 0`
18. checks if argument is a file / directory (in an IF): "`-f $arg`", or "`-d $arg`" (don't use "type", like in a for)
 - a. `-x` is for executables
 - b. man test shows all these types
19. `chmod u=rwx,g=rx,o=r myfile` (permissions for owner, group, others)
20. **elif** not else if
21. `[$foundFile -gt 0 -a $foundDirectory -gt 0];`
 - a. `-a` means `&&`
 - b. `-o` means `||`
22. when you want to remove an entire line: `sed "/pattern/d" $file`
23. be careful when using sed to check for entire words, so "`\<word\>`" not just "word"
24. - when you want to grep a string (which is not a file name): `echo $string | grep "pattern"`
25. you can do: `if grep "pattern" then ..`

26. Use printf syntax in awk: `awk '{printf "%s@scs.ubbcluj.ro", $1}'`
27. use the command "file" to determine the type of a file
28. for the sleep command we need `unistd.h`
29. `fgets(string, MAX_STRING_LENGTH, stdin)` ⇔ read a line from the keyboard
30. `srand(time(NULL))` ⇔ reset the random generator
31. `sed "s/\([0-9]\)/\1\1/g" file.txt` ⇔ prints a file with all the digits duplicated (if we also use `-i` => the file is actually modified)
32. `grep "[aeiouAEIOU]$" aux.txt` ⇔ prints the lines which end in vowels
33. `echo $(dirname `pwd`)` ⇔ prints the parent directory
 - a. `echo ~` ⇔ prints the root
34. `if ! [-f $arg -o -d $arg];` ⇔ checks if the argument is not a file nor a directory (beware of the negation sign)
35. A non-zero return value in shell => false; 0 => true
36. A child process is a zombie from the moment it finishes execution until its parent calls `wait` or `waitpid`
 - a. Call `wait` or `waitpid` for each process you create
37. CTRL-C ⇔ SIGINT ⇔ signal no. 2 ⇔ `kill -2 pid`
 - a. SIGTERM = 15
 - b. SIGKILL = 9
 - c. `man 7 signals` shows all of them
38. `signal(SIGINT, f)` => whenever SIGINT is received, the function `f` is called
 - a. This should be called at the beginning of the program
 - b. This works for all signals but SIGKILL - this cannot be stopped or modified
 - c. Needs `#include <signal.h>`
39. Whenever a child process stops, the parent receives a SIGCHLD
 - a. `signal(SIGCHLD, SIG_IGN)` will basically prevent zombie creation
40. SIGUSR1 and SIGUSR are "set aside" so that we can use them in any way we want
41. Running other programs:
 - a. Searches PATH for the program
 - i. **execvp**("grep", a)
 1. `char* a[] = {"grep", "/an1/gr911/", "/etc/passwd", NULL}`
 - ii. **execlp**("grep", "grep", "/an1/gr911/", "/etc/passwd", NULL)
 - b. Doesn't search PATH for the program
 - i. **execv**("/bin/grep", a);
 1. `char*a[] = {"/bin/grep", "/an1/gr911/", "/etc/passwd", NULL}`

- ii. **execl**("/bin/grep", "/bin/grep", "/an1/gr911/", "/etc/passwd", NULL)
- c. NULL marks the end of the arguments
- d. Path = all the places in which an executable might be located
- 42. First argument (0) of a program is always the command name
- 43. If we close a pipe before calling fork(), the child will inherit a closed pipe
- 44. pipe[0] = read; pipe[1] = write
- 45. **popen**:
 - a. We can call other programs/ scripts from a C program and also get its output or provide some input to it
 - b. Basically an exec which doesn't overwrite the current program
- 46. **dup/ dup2**
 - a. File descriptor table: basically every file in a program will have a handle, which indicates what we can do with it / knows how to operate it;
 - i. e.g. "open" will open a file to reading or writing -> its handle is stored in the fd table on the first available position; its index will be the value of the handle
 - ii. Positions 0, 1, 2 are usually "reserved" for (console) stdin, stdout and stderr
 - iii. A pipe will have 2 of these file descriptors: one for reading and one for writing
 - b. int **dup**(int oldFD) - copies the fd given as an argument to a new position in the table; so now, when we want to do a certain operation, we can do it from either of the fd's
 - c. int **dup2**(int oldFD, int newFD) - overwrites the old fd with the new one (so for example we can overwrite position 0 - stdin to have the handle for a fifo); it also **closes the old fd before overwriting it**
 - d. Exec will not overwrite the file descriptor table, so if we do a dup, then an exec, the new program will use the same fd table
 - e. The exec will also close the files that it used (so if we have a pipe, it will close the part that it used, but obviously not both ends)
 - f. **Undo** dup calls: before using the fd, make and store a copy of it using dup(), then do whatever you want with dup2(), then set the fd back to its previous value, which was stored in the beginning
- 47. Shared memory
 - a. IPC = interprocess communication
 - b. Each(IPC I think ?) has a unique ID; they remain in the system and we need to clean them up

48. Threads

- a. They don't copy data from the process (as opposed to fork), and will also need another stack
- b. If a thread fails, all other threads fail
- c. **Race condition**: multiple processes/ threads want to access the same resource (=critical resource) simultaneously
- d. **Semaphore** - doesn't allow more than n threads at a time to access a resource (n is set when initializing it)
- e. **Barrier** - waits for all n threads to reach it, and only then does it allow (all of) them to go on

49. Monolithic os

- a. they are basically one big program

50. Microkernel os

- a. it's like a dispatcher which transfers messages between different components; in this example, if a driver (which is one of the "components") fails, it does not bring the entire system down, like in the monolith's case
- b. The BIOS is kind of like a microkernel OS

51. Process states

- a. **HOLD** - a process which is almost ready to enter the system
- b. **READY** - the process is in the memory and is ready to be served by a processor
- c. **RUN** - the process is running
 - i. This will alternate quickly with

1. WAIT

- a. The process waits for a "slow" operation (usually disk operations); while it does that, it's no longer running, so that other processes can run

2. SWAP

- a. The RAM is full, so the OS "dumps" a processes' memory on the disk, temporarily, so that another process can run

- d. **FINISH** - the process is done

52. Semaphore = pair (v(s), c(s))

- a. v = the value of the semaphore, integer; when it's $\geq 0 \Rightarrow$ a new process can be run, otherwise it will be made to wait
- b. c = a queue which contains all the processes that are WAITing
- c. It has 2 operations

- i. **P** will try to run a process; if it has space ($v(s) > 0$) \Rightarrow we run it, otherwise it's pushed to the queue and the process is set to WAIT
 - ii. **V** will be called when ending a process; if there are processes in the queue (so if $v(s) < 0$), we will set one of them to READY and pop it out of the queue
 - iii. In a way, $P \sim \text{lock}$, $V \sim \text{unlock}$
- 53. Livelock - it's not a deadlock, bc things are moving; however, they are not progressing at all
- 54. Segmented memory - it defines a section of memory which will have certain permissions/ properties;
- 55. Loading a process
 - a. All at once - once it's loaded, it goes really fast; slow start-up time, might waste memory
 - b. Load each page as needed - kind of faster start-up, but the execution is slower, no wasted memory
 - c. "Locality principle" - usually when we load page x , we will also need pages $x+1$, $x+2$.. (ex. in a for loop)
- 56. **NRU** (not recently used) algorithm - for each page we store 2 bits; 1 of them knows whether the page has been referenced recently, the other if the page has been modified recently \Rightarrow 4 classes, where the most important is the one with both bits 1 \rightarrow modified but not referenced \rightarrow referenced but not modified \rightarrow both bits 0
 - a. Could be mixed with FIFO \Rightarrow better results
- 57. **LRU** (least recently used) - we have a matrix of size pages * pages; each time a page is used we fill its row with 1 and its column with 0 (in this order); the page with the least 1s on its row is the least used
- 58. When searching for where to allocate new, contiguous memory (malloc)
 - a. First fit - search for the first available area in which it fits
 - b. Best fit - search for an area s.t. the space left is the smallest possible; but we will be left with very small slices of memory
 - c. Worst fit - opposite of best fit
 - d. Buddy-system - we split the memory in chunks of size 2^k ; whenever we need to allocate x , we search for the first available power of 2 $\geq x$; if that is too large, we then divide what's left into other powers of 2, which will be free
- 59. **Caching**, from small and fast to the large and slow
 - a. Registers \rightarrow L1 \rightarrow L2 \rightarrow L.. \rightarrow RAM \rightarrow HDD/ SSD

- b. When we need to search for things/ their positions in cache =>
 - i. Direct caching ~ hashing, so we will have collisions (thrashing)
 - ii. No “hashing” - we just place it on the first available position, and when we need it, we search for it
 - iii. Compromise (set-associative caches) - we split the cache into larger sets which are accessed directly; but inside them we go linearly, like in the no hashing method
- 60. Symbolic link - a new, different name for the same inode
- 61. Hard link - a new inode for the same data; can only be created by the root